

# Visual Data Analysis with Task-based Recommendations

Leixian Shen\*  
Tsinghua University

Enya Shen†  
Tsinghua University

Zhiwei Tai‡  
Tsinghua University

Yihao Xu§  
Tsinghua University

Jianmin Wang¶  
Tsinghua University



Figure 1: User interface. When uploading a dataset (A), the data field shows data columns with the corresponding data type (B). Users can customize system settings, including interested data columns (C), max number of charts (C), recommendation mode (D), ranking scheme (E), and task list (F). Clicking *Recommendation* button will generate visualizations on the right. Selected tasks are tabbed, along with a *Display by task* switch (G). If the switch is open, the charts will be displayed by tasks. (H) is the default thumbnail view with a row of charts. Clicking the task tab will show all the recommendations (I). If the switch is closed, all the charts will be deduplicated first and displayed together like in Figure 8.

## ABSTRACT

General visualization recommendation systems typically make design decisions of the dataset automatically. However, most of them can only prune meaningless visualizations but fail to recommend targeted results. This paper contributes TaskVis, a task-oriented visualization recommendation system that allows users to select their tasks precisely on the interface. We first summarize a task base with 18 classical analytic tasks by a survey both in academia and industry. On this basis, we maintain a rule base, which extends empirical wisdom with our targeted modeling of the analytic tasks. Then, our rule-based approach enumerates all the candidate visualizations through answer set programming. After that, the generated charts can be ranked by four ranking schemes. Furthermore, we introduce combination recommendation, leveraging a set of visualizations to give a brief view of the dataset collaboratively. Finally, we evaluate TaskVis by a series of use cases and a user study.

\*e-mail: slx20@mails.tsinghua.edu.cn

†e-mail: sheneya@tsinghua.edu.cn

‡e-mail: tzw20@mails.tsinghua.edu.cn

§e-mail: yh-xu18@mails.tsinghua.edu.cn

¶e-mail: jimwang@tsinghua.edu.cn

## 1 INTRODUCTION

Data visualization is a promising way to facilitate data exploration. However, transforming the raw data into well-designed visualizations require users to be familiar with both the dataset and principles of effective visual encoding. Although the visual data analysis demand has nourished numerous visualization tools, they usually present a steep learning curve. Users need to engage in a tedious and time-consuming process to explore different design decision through trial and error. To address this issue, a remarkable series of Visualization Recommendation (VisRec) systems have been proposed both in academia [20, 25, 34, 37, 59, 61] and industry (e.g., Tableau and Power BI) to facilitate the visualization authoring process.

Over the past two decades, a surge of VisRec systems emerged as a powerful tool for visual analysis. The existing VisRec systems mainly consist of two categories [62]: rule-based approaches and learning-based approaches. Rule-based approaches build upon empirical visualization design knowledge. For instance, APT [35], Show me [36], DIVE [27], SeeDB [54], and Voyager [59] leverage predefined perceptual rules to automatically generate visualizations. These rule-based approaches are intuitive to understand and can guarantees good explainability of the recommendations. However, they suffer from high requirements for domain expertise. Learning-based approaches learn from plenty of well-designed visualization examples to train models for VisRec, such as DeepEye [34], Data2Vis [20], VizML [25], and Table2Charts [61]. Learning-based approaches avoid complicated hand-crafted rules, but they lack good scalability and may fail to achieve comparable performance as rule-




Figure 2: Recommendation pipeline. The recommendation engine accepts analytic tasks and data properties, and then automatically generates appropriate visualizations for users.

based methods due to the limited training data.

Although these VisRec systems can generate meaningful visualizations, they pay relatively little attention on the diversity of user intent. Previous studies have documented that the effectiveness of a visualization largely depends on the user’s analytic task [14, 41, 43]. For a dataset, different people may hold different expectation of the recommendations according to their analysis intent. However, most of the existing VisRec systems lack detailed modeling of analytic tasks, so they can only prune meaningless results but fail to recommend targeted visualizations. As a result, given specific data columns, they will always produce the same charts for all the users. Even if some works have taken analytic tasks into account, the comprehensiveness and depth are still insufficient. For instance, Draco [37] only groups tasks into two categories: value tasks and summary tasks. NL4DV [38] only integrates five types of low-level analytic tasks into VisRec system. Fortunately, there is a remarkable series of empirical wisdom on task modeling both in academia and industry. For example, academic work from Saket *et al.* [43] comprehensively evaluates the effectiveness of five commonly used chart types across different tasks. For industry, Financial Times Visual Vocabulary [11] provides great guidelines for designers and journalists to select the optimal symbology for data visualisations based on their tasks. However, all these design knowledge of analytic task are scattered across heterogeneous sources (e.g., papers, training sessions, projects etc.) and are rarely implemented in practice to power VisRec systems.

In response to the problem, we first summarize the empirical knowledge of the analytic task and then propose a task-oriented visualization recommendation system, TaskVis<sup>1</sup>. As shown in Figure 2, the recommendation engine accepts analytic tasks and data properties and automatically generates appropriate visualizations for users. TaskVis differs from existing VisRec systems by modeling analytic tasks in details. The contributions of this paper are summarized as follows:

- A task base of 18 common analytic tasks with their appropriate chart types by a survey both in academia and industry.
- A rule base, which extends empirical wisdom with our targeted modeling of analytic tasks.
- A task-oriented VisRec system, TaskVis, which allows users to select their tasks precisely.
- A series of use cases and a user study to prove the effectiveness of TaskVis.

## 2 RELATED WORKS

TaskVis draws upon prior efforts in analytic task construction and visualization recommendation systems.

### 2.1 Analytic Tasks

There is a growing body of literature recognizes that the analytic task is vital for visualization design. Vartak *et al.* [53] proposed that the intended task is one of the five factors that must be accounted for while making visualization recommendations. There

have been numerous task taxonomies in the data visualization literature providing definitions of analytic tasks. For instance, Amar *et al.* [13] presented a set of ten low-level analytic tasks that capture the user’s activities while employing visualization tools for data exploration, including *retrieve value*, *filter*, *compute derived value*, *find extremum*, *sort*, *determine range*, *characterize distribution*, *find anomalies*, *cluster*, and *correlate*. Saket *et al.* [43] evaluated the effectiveness of five commonly used chart types (mark as table, line, bar, scatter, and pie) across the ten tasks [13] by a crowdsourced experiment. Kim *et al.* [31] measured subject performance across task types derived from the ten tasks mentioned above, and included *compare values* task in addition. The tasks were further grouped into two categories: value tasks that just retrieve or compare individual values and summary tasks that require identification or comparison of data aggregation. Draco [37] and Dziban [33] are the following VisRec systems that involve value tasks and summary tasks. NL4DV [38] additionally included a *Trend* task. AutoBrief [30] further enhanced VisRec systems by introducing domain level tasks. Wang *et al.* [56] summarized 11 categories of task that are commonly adopted in fact sheets. Deep into scatter charts, Sarikaya *et al.* [44] first collected model tasks from a variety of sources to formulate the seeds for scatterplot-specific analytic task list, then performed card-sort to group tasks together based on their similarity. Rather than inputting task proactively, behavior-based systems represented by HARVEST [23] and Eye Gaze [51] capture users’ interaction behavior to infer their analytic task. Besides, various empirical materials aiming at improving chart literacy have gradually formed in the industry practice. For instance, choosing a good chart [4] identified four tasks for market researchers and Financial Times Visual Vocabulary [11] summarized nine tasks commonly used by designers and journalists. However, in the visualization community, there lack a comprehensive summary and further application of these tasks.

### 2.2 Visualization Recommendation

Early VisRec works represented by APT [35], SAGE [42], and Rank-by-Feature Framework [46] are mostly rule-based. They take data features (e.g., statistical properties) into account and leverage visualization rules to prune large search space. The design rules in the community increase iteratively as subsequent works embody the previous works and produce new rules. Voyager [58, 59], DIVE [27], Show me [36], Polaris [52], Profiler [29] all contributed to the enrichment of rules with more data types. Instead of considering common visualization types, Wang *et al.* [55] developed rules for automatic selection of line graph or scatter plot to visualize trends in time series. Datesite [17] proactively generated insights using automatic algorithms for data exploration. To better benefit from design guidance provided by empirical studies, Moritz *et al.* [37] proposed a formal framework that models visualization design knowledge as a collection of ASP (Answer Set Programming) constraints [22]. We also merge this idea into TaskVis and make further improvements.

Rule-based approaches highly depend on domain expertise. In response, learning-based systems draw from existing well-designed visualizations to train neural models. DeepEye [34] trained a decision-tree-based binary classifier to determine whether a visualization is good or not, and a learning-to-rank model to rank visualizations. Data2Vis [20] built a multi-layered attention-based encoder-decoder network to directly map JSON-encoded datasets to Vega-Lite specification. Draco-Learn trained a RankSVM model on ranked pairs of visualizations and then automatically learned weights for soft constraints in ASP. VizML [25] identified five key design choices while creating visualizations. For a new dataset, the 841 dataset-level features extracted were fed into neural network models to predict the design choices. Table2Charts [61] used a deep Q-learning approach with copying mechanism and heuristic searching. Qian *et al.* [39] proposed an end-to-end machine learning approach while MultiVision [60] built a Siamese neural network.

<sup>1</sup> <https://github.com/ShenLeixian/TaskVis>



Figure 3: Example of Vega-Lite specification. The figure shows a bar chart with an ordinal variable at x axis, a quantitative variable at y axis, and a nominal variable at color channel. In addition, sort, sum aggregation, and zero stack transformation are applied.

### 3 BACKGROUND

#### 3.1 Vega-Lite

Inspired by prior works [33, 37], we leverage Vega-Lite to specify visualizations. Vega-Lite [45] is a high-level language of interactive graphics. It provides a composition declarative JSON syntax for visualization displays. Vega-Lite supports various atomic components, including marks, channels, and transforms. Figure 3 shows an example of Vega-Lite specification, which is a bar chart, including x-axis, y-axis, and color channel. In addition, sort, sum aggregation, and zero stack transformation are applied.

#### 3.2 Answer Set Programming

Answer Set Programming (ASP) is a declarative programming paradigm based on logic programs and their answer sets. It provides a simple yet powerful modeling language to solve combinatorial problems. An ASP program is built on *atoms*, *literals*, and *rules*. *Atoms* are basic building blocks in ASP to represent elementary propositions. *Literals* are atoms *A* or their negation *not A*. A *rule* in ASP is in the form of  $A :- L_0, L_1, \dots, L_n$ , where  $L_i$  is a literal. The head (*A*, left of  $:-$ ) in the rule is true only if all the literals in the body ( $L_i$ , right of  $:-$ ) are true. For instance,  $\text{road}(\text{cityA}, \text{cityC}) :- \text{road}(\text{cityA}, \text{cityB}), \text{road}(\text{cityB}, \text{cityC})$ , states that only when city A and B are connected and city B and C are connected, city A and C can be connected. If a rule is headless, it means that satisfying all the literals in the body will lead to a contradiction. Illustrated with example,  $:- \text{on}(\text{computer}), \text{off}(\text{computer})$ , states that a computer cannot be both on and off. If a rule is bodyless, it asserts the fact that its head is true, for example,  $\text{on}(\text{computer})$ , states that the computer is on. Besides, the aggregate rule  $p \{A_0, \dots, A_n\} q$ , restrains that at least *p* and at most *q* atoms in the set are true. More modeling constructs of ASP can be found in its document<sup>2</sup>. Clingo [22] is an ASP system to solve logic programs. Given a combinatorial problem, a series of rules make up the ASP program, Clingo can be used as a solver to quickly enumerate all eligible combinations, with which the researchers can concentrate on the actual problem.

### 4 OVERVIEW

**User interface:** As shown in Figure 1, when uploading a dataset (A), the data field shows data columns with the corresponding data type (B). Users can customize system settings, including interested data columns (B), max number of charts (C), recommendation mode (D), ranking scheme (E), and task list (F). Clicking *Recommendation* button will generate visualizations on the right. Selected tasks are tabbed, along with a *Display by task* switch (G). If the switch is open, the charts will be displayed by tasks. (H) is the default thumbnail view with a row of charts. Clicking the task tab will show all the recommendations (I). The chart can be previewed by clicking on it like in Figure 6. If the switch is closed, all the charts will

be deduplicated first and displayed together like in Figure 8. The combination recommendation view is shown in Figure 9.

**Architecture:** The architecture of TaskVis is shown in Figure 4. Task base (a) is the foundation of the system, which will be discussed in Section 5. Apart from the prerequisite dataset, users can also input optional data columns of interest and analytic tasks (b) like in Figure 1(B, F). The preprocessing module (c) will later extract data features (e.g., data type and cardinality) automatically. User intervention is also allowed in this phase as data type may be ambiguous in some scenarios. After preprocessing, extracted data features together with the user’s input are sent to visualization generation module (d). Conversion to ASP constraints is a premise to make input available for Clingo, which is an ASP system to solve logic programs. Apart from the user’s input, a rule base (e) is the core of TaskVis and also should be sent into Clingo. More details will be discussed in Section 6.2. All enumerated results in ASP constraints are transformed into Vega-Lite specification at the end of visualization generation stage. In visualization ranking module (f), if previous modules have processed multi-tasks, the output VisRec list can be deduplicated. After that, each Vega-Lite is assigned a *cost score* based on a cost model inspired by GraphScape [32], which will be discussed in Section 7.1. On the basis, four rank schemes are designed for users to select, and each scheme has applicable scenarios, which will be discussed in Section 7.2. Finally, a VisRec list (g) is presented to users, where each chart in the list is an independent recommendation. So far, we have presented a complete visualization recommendation workflow. Iterating multiple tasks, we can obtain numerous visualizations that describe partial dataset. In addition, instead of just presenting individual charts, we propose combination recommendation (h) to extend the system, which leverages a set of charts to collaboratively describe the dataset based on task-oriented design principles. More details will be discussed in Section 8.

### 5 ANALYTIC TASKS

While many task taxonomies have been constructed, most of them only apply to specific scenarios. Besides, based on existing task taxonomies, some new tasks introduced by incremental works for visual data analysis are scattered across heterogeneous sources. To summarize these tasks, we maintain a task base from three aspects:

- **Empirical academic studies:** A large body of early works pay attention to the effectiveness of visualization types for a selected task, either special for a particular chart type (e.g., scatter [44, 55], bar [16, 21], circle [16, 21], error bar [15], arc [49], area [49], and line [55]) or comprehensive evaluation [13, 18, 24, 31, 43, 57].
- **Empirical summaries in industry:** Various fields such as newspaper, finance, geography, and engineering have a strong requirement for data visualization, rich practice experience has also been accumulated in the form of guides, training sessions, projects, libraries, etc. [2–4, 6, 9–11]
- **Customized tasks:** Apart from commonly used tasks, task base also contains some others that we consider meaningful during our visualization practice, such as *error range*.

Based on the studies above, similar to the methodology to identify scatterplot tasks in [44], we first formulate the seeds for analytic task list, covering both low-level and high-level tasks. To abstract these analytic tasks, we then invite several data visualization researchers with more than five years of experience to perform a card-sort and merge the similar tasks. Finally, we summarize 18 classical analytic tasks and their appropriate chart types as shown in Table 1. What needs to explain is that *filter* task requires users to specify some constraints of data additionally. So the implementation of this task is placed in data preprocessing stage, allowing users to filter data according to their own needs. For the other 17 tasks, users can choose anyone to perform data analysis at will. We also extract appropriate chart types for each task and retained the priority. For instance, the most appropriate chart type for *comparison* task is line

<sup>2</sup><https://potassco.org/>



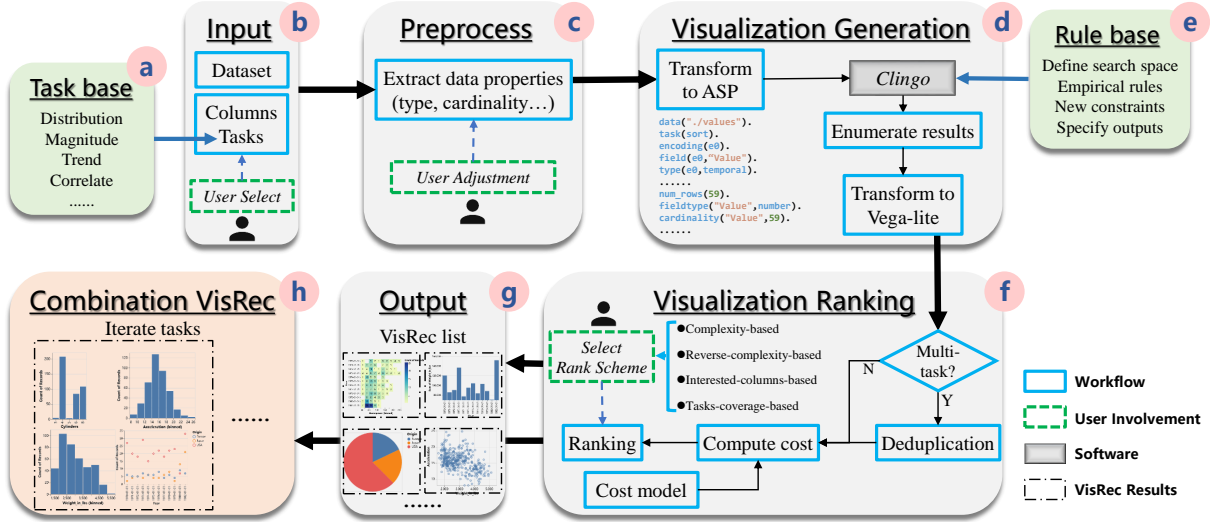


Figure 4: Architecture of TaskVis. TaskVis consists of six modules and two bases (for tasks (a) and rules (e) respectively). (b) Input: accepts the user's input. (c) Preprocess: extract data properties. (d) Visualization Generation: enumerate all qualified candidates. (f) Visualization Ranking: rank all visualizations according to selected scheme. (g) Output: present recommendation results to users. (h) Combination VisRec: make combination recommendations by iterating tasks.

chart, followed by scatter chart, and finally bar chart. (\*) in the table indicates combination of marks, for example, rect(text) represents that a text layer is superimposed on rect chart.

## 6 VISUALIZATION GENERATION

Generally, VisRec systems need to enumerate all possible visualizations first and then make recommendations [40].

### 6.1 Visualization components

Vega-Lite includes various atomic components, we support a part of them as follows:

**Mark:** To better assist data analysts, we enrich the chart types of recommendations as much as possible. To this end, we merge and filter all the marks support in Vega-Lite and finally select 14 types, arc, area, bar, box plot, circle, error band, error bar, geoshape, line, point, rect, rule, text, and tick. Besides, TaskVis also supports layered charts (e.g., line + point) in some tasks. Compared with existing VisRec systems, TaskVis can recommend richer charts.

**Channel:** In visual analysis, the rational application of visualization channels is one of the pivotal factors for designing aesthetic visualizations. Visualization channels include location, color, size, shape, slope, texture, and animation. There is a consensus that the first four can generally meet the demands of visualization design. Hence for simplicity, encoding channels in TaskVis only refer to the axis (x and y) and legend (color, size, and shape). It is worth mentioning that at most two legends are involved in a chart.

**Transform:** Data transformation enables new value generation by various operations. For example, we can use *bin* to bucket quantitative or temporal data, discretizing continuous variables. *Regression* and *loess* can be applied to produce trend lines. All the transform operations of Vega-Lite included in TaskVis are bin, aggregate, sort, stack, regression, loess, and scale.

### 6.2 Rule base

Various visualization components form a huge search space, among which there are numerous invalid combinations and meaningless visualizations that show limited information. Fortunately, there are numerous rules either from users or traditional wisdom to prune "bad" visualizations. However, if all these rules are implemented

in code with branch structure, it will be complicated to manage. Inspired by Draco [37], we leverage answer set programming to construct design knowledge in a unified and extensible manner and finally build a rule base, which consists of three sources: empirical wisdom, customized rules for task, and user partial specification.

#### 6.2.1 Empirical wisdom

Draco [37] maintains a knowledge base in ASP depending on prior efforts such as APT [35] and CompassQL [58]. ASP program of the knowledge base includes visualization attributes declaration, preference over design space, and output specification. Visualization attributes declaration specifies the domains of attributes, for example, `type (quantitative;ordinal;nominal;temporal)`, defines the four high-level data types and `channel (x;y;color;size;shape)`, defines the five supported single encoding channels. Preference rules over design space can effectively prune the search space, for example, `:- channel(E, shape), not type(E, nominal)`, ensures that the model will not consider applying shape channel on non-nominal variables. Specification of outputs like `#show channel/2` define the output items and the number of its parameters. With the aforementioned modeling process, Draco can be leveraged to build increasingly sophisticated visualization systems with less efforts. In addition to Draco, we also integrate appropriate rules from other sources as listed in Table 1.

#### 6.2.2 Customized rules for task

Draco [37] only includes two abstract tasks (value and summary), which can hardly capture the user's activities during visual analysis. In response, TaskVis extends Draco with additional detailed modeling of analytic tasks based on the task base in Section 5. TaskVis introduces an additional primitive `task()`. As shown in Table 1, we summarize appropriate chart types for each task. It can be formulated as rules to effectively narrow the search space, for example, `:- task(sort), not mark(bar)`, is an integrity constraint stating that it cannot be the case that mark type applied in *sort* task is not bar. There is another type of rule that restricts operations to several specific tasks. For example, aggregation is not allowed in *determine range*, *trend*, *deviation*, and *error range* task, as they should display the raw data. Besides, based on the characteristics of each task, we also formulate separate rules independently. For instance, `:- channel(E, x), not type(E, temporal)`, `task(change-over-time)`, states that the variable

Table 1: Task base. *Mark* column lists appropriate marks, where the rank has priority, (\*) indicates the combination of marks, e.g. rect(text) means a text layer is superimposed on rect chart.

Task	Mark	Description	Reference
Change Over Time	line/area	Analyse how the data changes over time series	[3, 6, 9–11, 24]
Characterize Distribution	bar/point	Characterize the distribution of the data over the set	[2–4, 6, 11, 13, 18, 38, 43, 44, 50, 56]
Cluster	bar/point	Find clusters of similar attribute values	[3, 13, 43, 48, 50]
Comparison	line/point/bar	Give emphasis to comparison on different entities	[2, 4, 6, 9, 10, 26, 27, 31, 44, 56]
Compute Derived Value	rect(text)/arc/bar	Compute aggregated or binned numeric derived value	[13, 27, 43, 50, 56]
Correlate	bar/line	Determine useful relationships between the columns	[2–4, 6, 9, 11, 13, 18, 24, 27, 38, 43, 44, 48, 50, 56]
Determine Range	tick/boxplot	Find the span of values within the set	[6, 13, 43, 50]
Deviation	bar(rule)/point(rule)	Compare data with certain value like zero or mean	[11]
Error Range	errorband/errorbar	Summarizes an error range of quantitative values	[15]
Filter	rect/bar/arc	Find data cases satisfying the given constrains	[13, 38, 43, 50, 56]
Find Anomalies	bar/point	Identify any anomalies within the dataset	[13, 18, 24, 26, 43, 44, 48, 50, 56]
Find Extremum	bar/point	Find extreme values of data column	[13, 26, 31, 43, 50, 56]
Magnitude	arc/bar	Show relative or absolute size comparisons	[6, 11, 49]
Part to Whole	arc	Show component elements of a single entity	[2, 4, 6, 9, 11, 49, 56]
Retrieve Value	rect(text)	Find values of specific columns	[3, 10, 13, 26, 31, 38, 43, 50, 56]
Sort	bar	Rank data according to some ordinal metric	[11, 13, 24, 43, 50, 56]
Spatial	geoshape/circle(text)	Show spatial data like latitude and longitude	[3, 6, 9–11]
Trend	point(line)	Use regression or loess to show the variation trend	[2, 27, 55, 56]

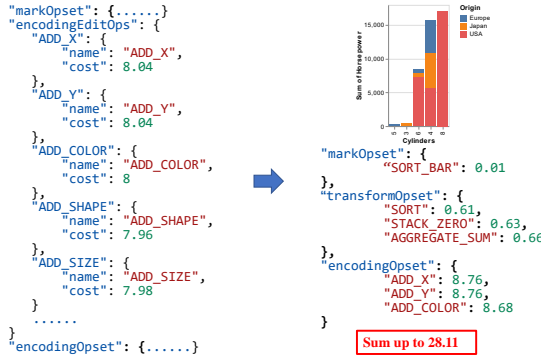


Figure 5: Example of cost model. *cost score* of the visualization in Figure 3 can be obtained by summing up cost of all components.

on x-axis in *change over time* task must be temporal. We develop these rules through iterative informal feedback from researchers and students with varying-level experience with visualizations based on a prototype of TaskVis. In each round, we pick out meaningless charts and further improve our rule base.

### 6.2.3 User partial specification

TaskVis supports full-automatic recommendation, users only need to input the dataset and the system can automatically generate appropriate visualizations. However, in order to make user input more flexible, we allow inputting partial specifications which includes interested data columns and their analytic tasks. It helps to restrict the search space within well-formed specifications of the user’s interests. These rules are generated based on interactions between users and visualization interface, as shown in Figure 1(B, F).

Recall Figure 4 (d), after transforming to ASP programs, the user’s input along with the rule base is sent into Clingo. Differ from Draco [37], which eliminates meaningless visualizations using hard constraints and finds the most preferred charts according to soft constraints. TaskVis first enumerates all candidate visualizations in ASP by hard constraints, and then converts to Vega-Lite specification for visualization ranking module to handle. Each selected task corresponds to an independent VisRec list.

## 7 VISUALIZATION RANKING

TaskVis allows users to input multiple tasks, and some tasks have overlapping recommendations. This means that one chart may be applicable to multiple tasks. Therefore, if the visualization generation module handles multi-tasks, the VisRec list can be deduplicated first. After that, we derive a cost model inspired by GraphScape [32] to quantify the visualizations and design four rank schemes.

### 7.1 Cost model

As TaskVis is powered by Vega-Lite, a single visualization can be split into independent components, as shown in Figure 3. GraphScape [32] introduces a graph structure for modeling relationships among Vega-Lite-based charts. Graph nodes represent visualizations and edges represent edits that transform one visualization to another. GraphScape leverages linear programs to derive transition costs via a partial ordering of visualization components. Illustrated with example, in the transition category of visual encodings, following the perceptual kernels method [19], metric-space embeddings are performed on visual encoding type comparisons. The results find the following transition cost order:  $x, y > \text{row, column} > \text{color} > \text{size} > \text{shape} > \text{text}$ , which assumes that changes to spatial encodings require more interpretation effort. The corresponding generated cost is shown in Figure 5 (left).

Inspired by GraphScape [32], we also generate a cost model and define *cost score* to measure the complexity of different components. For a chart, *cost score* can be obtained by adding the corresponding cost values of all components. Lower *cost score* means the visualization is more concise. In detail, we extend GraphScape with the supplement and subdivision of visualization components. For mark type, we refer Table 1 to generate cost with priority of each task. Among channels, we replenish some components that are special for arc and geoshape chart. As to data transform, we not only add operations like stack, regression, and loess, but also subdivide aggregate operation into AGGREGATE.SUM, AGGREGATE.COUNT, and AGGREGATE.MEAN. Based on the cost model, each visualization can be assigned a *cost score* with good explainability. For example, in Figure 3, the chart is generated under *sort* task, its component set and corresponding cost are shown in Figure 5 (right), *cost score* can be obtained by summing up cost of all components.

## 7.2 Rank schemes

After computing *cost score* of each chart, four rank schemes are designed as follows:

**(I) Complexity-based ranking:** This scheme ranks visualizations according to the chart's *cost score* from low to high. When people make visualization observations, they tend to view charts as a whole and hope to understand the information inside charts in a short time. According to our cost model, charts with lower *cost score* are easier to understand, which is in line with people's habit of observing charts. This scheme is appropriate for most analytic tasks, such as *change over time*, *characterize distribution*, *cluster*, etc.

**(II) Reverse-complexity-based ranking:** Overall, charts with higher *cost score* come first in this scheme. However, instead of reversing the visualizations in complexity-based ranking scheme, we first adopt the unsupervised algorithm DBSCAN to cluster all the results into several categories, where visualizations in one category are similar to each other (e.g., just exchange the coordinate axis or replace an aggregation function). Then the visualizations are ranked by category in reverse order while maintaining the partial rankings within the category. The scheme is appropriate for tasks like *sort*, and *determine range*. The rationality is that charts of these tasks are relatively simple, charts with high *cost score* can present more information within a range that humans can easily perceive.

**(III) Interested-data-columns-based ranking:** This scheme is mainly related to the data columns of interest inputted by the user. Users usually expect a chart to display as many data columns of interest as possible. So if a chart fails to show all the user's interested columns, it will be assigned a penalty. The new *cost score* is computed by dividing ( $N1/N2$ ), where  $N1$  is the number of data columns covered in a single chart and  $N2$  is the number of the user's interested data columns. Visualizations are finally sorted by new generated *cost score* from low to high. The scheme fits some comprehensive tasks like *magnitude* and *find anomalies*.

**(IV) Tasks-coverage-based ranking:** The above three rank schemes are more suitable for single-task situations. For multiple tasks, we believe that if a chart is an overlap of multiple tasks' VisRec list, it must have universal meaning and should be ranked in the forefront. We define *task coverage number* of a chart as the number of tasks (within the user's input) it can satisfy. Before ranking by this scheme, we merge and deduplicate all tasks' VisRec list first, and identical visualizations average the *cost score* value. On the basis, all charts are first ranked by *task coverage number*, where those with the same *task coverage number* are partially ranked by the *cost score*. Noting that this rank scheme is valid only when TaskVis accepts multiple tasks, otherwise it will degenerate into complexity-based ranking scheme.

Back to Figure 4 (f), users are allowed to select the four rank schemes freely to deal with the generated visualizations. We also recommend appropriate rank scheme for each task as default in the system based on the materials mentioned in Section 5. In the output VisRec list (g), each chart is an independent result.

## 8 COMBINATION RECOMMENDATION

In this section, instead of recommending individual charts, we focus on generating a set of visualizations to give a brief view of the whole dataset, as shown in Figure 4(h). We call this issue *combination recommendation*. This is the gospel of users who are not familiar with the data and it is an effective way of system cold start. To this end, we extend tasks-coverage-based ranking scheme in the following two situations:

**Without Task selection:** If users select no tasks, the generation of candidate visualizations relies on iterating all the tasks in task base. Combination recommendations are further formulated according to the following principles: (1) The set of charts must cover all the data columns of interest to users. (2) Larger *task coverage number* is

always chosen first. We describe the combination recommendation algorithm in Algorithm 1.

**With Task selection:** If users have selected tasks of interest, iterating the inputted tasks can generate candidate visualizations. The first principle mentioned above will be relaxed here, because some data columns are only applicable for specific tasks, such as latitude and longitude can only appear in *spatial* task. If only a single task is accepted, TaskVis can also recommend a set of targeted charts to describe the task.

---

### Algorithm 1 Combination Recommendation

---

**Require:** *Result*: Visualization class; *Result.Fields*: Columns involved in a chart; *Cols<sub>interest</sub>*: Columns of interest that users input; *Recs*: *list[Result]*: Visualization list after deduplication.  
**Ensure:** A set of visualizations *ANS*: *list[Result]*;  
1: sort *Recs* by Tasks-coverage-based ranking scheme;  
2: set *Ans* = *Cols<sub>covered</sub>* = [];  
3: **while** *Cols<sub>covered</sub>*  $\subset$  *Cols<sub>interest</sub>* and *Recs* != *NULL* **do**  
4:   append *Recs*[0] to *ANS*;  
5:   *chosenfields* = *Recs*[0].*Fields*;  
6:   *Cols<sub>covered</sub>* = *Cols<sub>covered</sub>*  $\cup$  *chosenfields*;  
7:   **if** *chosenfields.length* is 1 **then**  
8:     *enums* = all data column combinations in *chosenfields*;  
9:   **else**  
10:     *enums* = all data column combinations in *Cols<sub>covered</sub>*;  
11:   **end if**  
12:   **for** *vis* in *Recs* **do**  
13:     **if** *vis.Fields* in *enums* **then**  
14:       delete *vis* in *Recs*;  
15:     **end if**  
16:   **end for**  
17: **end while**

---

## 9 EVALUATION

This section will discuss the evaluation of TaskVis by a series of use cases and a user study, along with comparison with existing technologies.

### 9.1 Use Cases

This section will illustrate the recommendation capability of TaskVis with use cases. The individual recommendation will be exhibited first, and the combination recommendation follows. More recommendations can be found in the supplemental materials.

#### 9.1.1 Individual Recommendation

Imagine that we are exploring a COVID-19 dataset of U.S. [5], which contains the number of daily confirmed and dead cases. We may be interested in the geographical distribution of deaths. The recommendations in the *spatial* task are shown in Figure 6. One chart is clicked for preview, it applies color channel on the circle mark, and the concrete numbers are labeled. Geoshape charts leverage heat maps to reveal the distribution of deaths in each state. For this dataset, we may also be keen on "how the number of daily confirmed cases has changed since the epidemic?". After selecting the *change over time* task and data columns of interest, the recommendation results are shown in Figure 7. We can intuitively observe the change in a total number of daily confirmed (national and four major regions distributed) by line and area charts, and the area charts also show the stack of data.

For the Happiness Ranking dataset [12], after choosing *Region*, *Happiness Score*, and *Health* columns, selecting six tasks, and closing the *Display by task* switch, all the charts are deduplicated and displayed together. As shown in Figure 8, the tasks-coverage-based ranking scheme is applied, charts ranked in the forefront are appropriate for more tasks, and the names are placed above the chart.

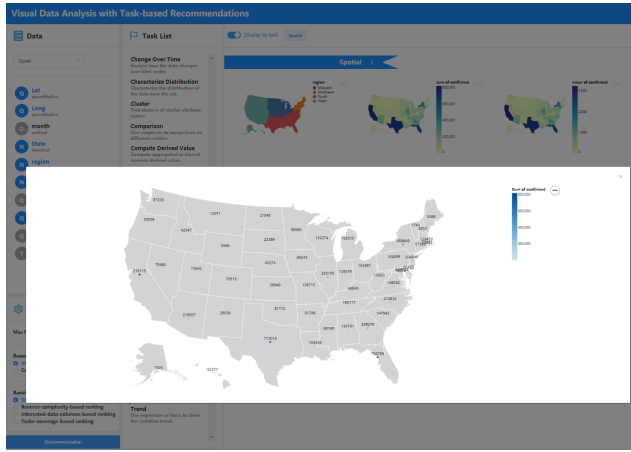


Figure 6: VisRec examples of the COVID-19 dataset in the *spatial* task with multiple columns. One chart is previewed by clicking on it.

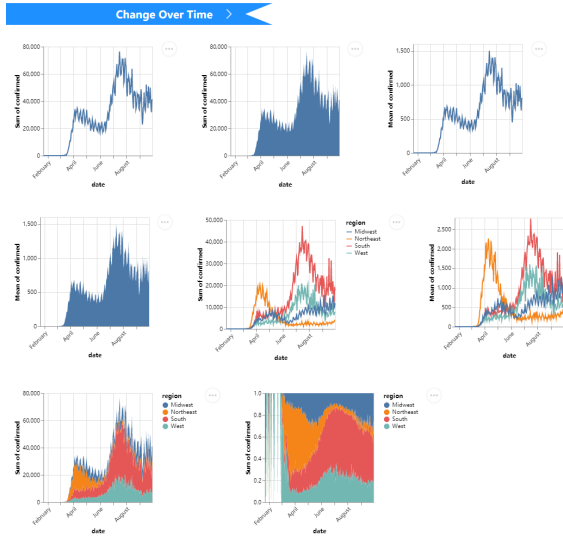


Figure 7: VisRec examples of COVID-19 dataset in the *change over time* task with the *region*, *confirmed*, and *date* columns.

The charts also support simple interactions like hovering to see data details.

### 9.1.2 Combination Recommendation

Imagine that we are analyzing the Hollywood Stories dataset [7] and are not familiar with the data. A general view of the dataset can be helpful for our subsequent in-depth exploration. As shown in Figure 9, without selecting tasks and data columns, the set of charts covers all the data columns and satisfies the most tasks. As the system displays thumbnails by default, high-cardinality variables (like *name*) will make the chart crowded, but clicking on the chart like in Figure 6 can view the original clear chart. Although the use cases present promising results, our target is not to show that combination recommendation results are always superior. We argue that it offers an option to obtain a general view of the dataset.

## 9.2 User Study

The overall goal of the study was to investigate TaskVis’s ability to recommend targeted charts for the user’s analytic tasks.

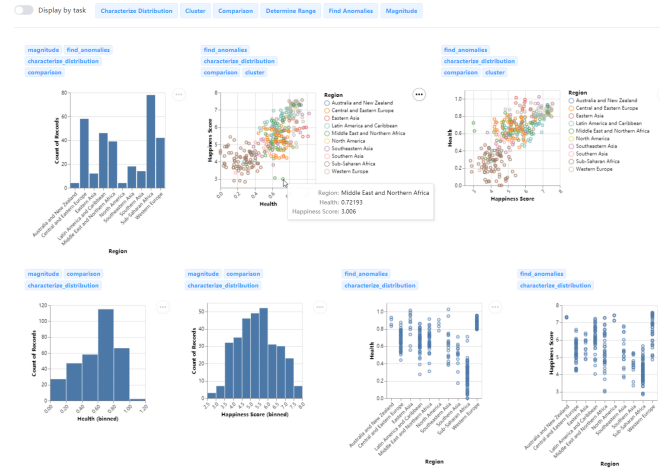


Figure 8: VisRec examples of the Happiness Ranking dataset. The *Display by task* switch is closed.



Figure 9: Combination Recommendation results of the Hollywood Stories dataset.

### 9.2.1 Participants

We invited 22 persons (6 female and 16 male) to participate the study. The participants include 1 undergraduate, 13 master candidates, 6 phd candidates, 1 researcher, and 1 data analyst. Among them, 8 participants perform data analysis every day, 9 monthly, and 5 less than once a month. All of them have the experience of using Microsoft Excel and programming with python to analyze and present their data.

### 9.2.2 Study Protocol

The study can be roughly classified into three parts: pre-designed experiment, free exploration, and post-study interviews. Before the study, we presented a short examination on color blindness and basic visualization knowledge. In the pre-designed experiment, we first made an explanation of the process in details and gave a brief description of each task. Then, given a recommended chart, participants were asked to answer a 5-point Likert scale question (-2 = strongly disagree, -1, 0, 1, 2 = strongly agree), “The chart of {1} dataset can satisfy the {2} task.”, where {1} is the dataset name and {2} is the task name. Metadata about the chart (e.g., task description and data columns) is also presented to help participants make their decisions. In total, there were 287 visualizations for participants to

cope with. After the experiment, we made a tutorial to introduce the system interface and show some examples. Participants were later asked to freely explore TaskVis through a web browser on their computer, especially to perform tasks related to their actual day-to-day work. The user study lasted about one hour. Participants did not receive any intervention during the whole process. Finally, we conducted post-study interviews to get participants' feedback, including a questionnaire guided by the system usability scale to gauge the usability of TaskVis.

### 9.2.3 Datasets

We provided participants with two datasets in the pre-designed experiment. One is the COVID-19 data of the United States [5], including the daily numbers of confirmed cases and deaths in each state. The dataset has 12903 records and consists of 5 quantitative, 3 nominal, and 1 temporal columns. The other is sales data of 406 different cars [1], which is composed of 5 quantitative, 2 nominal, 1 ordinal, and 1 temporal columns. For free exploration, we offered more datasets covering areas of weather, sport, movie, etc.

### 9.2.4 Analysis of result

All participants finished the pre-designed experiment and explored the system with several datasets. We finally collected 6249 valid records from 22 participants in the pre-designed experiment. The average score is 1.05. In general, TaskVis can recommend targeted charts to satisfy the user's analytic task. We further analyzed the charts that were scored -2 or -1 more than three times and summarized the reasons as follows: (1) The overlap between marks makes charts difficult to distinguish, especially when a large amount of data is presented; (2) Outliers affect the users' understanding of the chart. For example, in the COVID-19 dataset, the number of deaths in one day exceeds 5000, which is much higher than the average. This may lead to some charts presenting in extreme ways; (3) The chart conveys limited information. This is mainly reflected in the *determine range* task, where the chart only presents discrete variables. We think that all the reasons above are mainly aroused by the dataset, while we pay more attention to the form in which the data is presented. On the other hand, the abnormality of charts can also mirror the abnormality of the dataset, which is also meaningful in data analysis. So we believe that TaskVis still has a good recommendation ability, and we will also handle the problem in future work.

TaskVis was generally considered as usable. Participants mostly agreed with the statements "I think that I would like to use this system frequently." ( $\mu = 4.3$  on a 5-point Likert scale) and "I thought the system was easy to use." ( $\mu = 4.5$  on a 5-point Likert scale). The mean system usability score is  $\mu = 71$ , being between "OK" and "good" [8]. We also collected many feedbacks of the system. P4 (female, 27) said that "The integrated tasks cover commonly used tasks in my daily work and it (TaskVis) can help me quickly generate appropriate visualizations." For ranking effectiveness, P1 (male, 23) made a comment that "The four ranking schemes are interesting. I can choose different schemes according to my application scenarios. Usually, the charts listed in the front are relatively more appropriate." P18 (female, 23) commented regarding the final presentation of the system that "multiple presentation modes allow flexible exploration." In addition, P9 (male, 23) said that "Combination recommendation is useful, sometimes I select multiple tasks, but individual visualization may fail to cover all the tasks." Some participants also pointed out the drawbacks of TaskVis. P13 (male, 25) said that "It would be better to make the output editable." and P14 (male, 23) stated that "More data insights can be further integrated. More convenient interaction paradigms may enhance the overall user experience, such as pen, touch, and natural language." Overall, participants all agreed that TaskVis is an effective visualization tool to facilitate data exploration.

## 9.3 Comparison

TaskVis is a task-oriented VisRec system. There is no ideal technology in the community to compare with. So this subsection only provides qualitative discussions. To our knowledge, most existing VisRec systems can not accept the user's analytic tasks and fail to recommend targeted results. Instead, TaskVis pays more attention to the user's tasks and can generate richer charts. So the recommended visualizations of TaskVis can better satisfy the user's diverse analysis intents. Besides, most VisRec systems generate only a few charts in a recommendation circle. Instead, TaskVis can recommend various visualizations for users to choose, and the charts ranked in the front are relatively more appropriate. Different schemes also apply to different scenarios.

Recently, Visualization-oriented Natural Language Interfaces (V-NLIs) are becoming a complementary input modality to traditional WIMP interaction [47]. V-NLIs can extract the user's task by understanding the NL utterances since they may hint at the user's analysis goals. However, most V-NLIs can only deal with well-structured queries. They still face various challenges, such as the ambiguous and underspecified nature of human language. Compared with these technologies, TaskVis provides a set of commonly used analytic tasks and design an interaction to allow users to select their tasks precisely. Besides, the task base can be extended with more tasks in the future, and the system can easily integrate new tasks. Finally, we argue that TaskVis is not always superior to the aforementioned technologies. Rather, it provides a new approach to integrate the user's analytic tasks into VisRec systems and builds a task-oriented system that can generate targeted visualizations. We believe that these technologies and TaskVis can complement each other.

## 10 CONCLUSION AND FUTURE WORK

In this paper, we present a task-oriented visualization recommendation approach called TaskVis. TaskVis improves the existing technologies by strongly correlating recommendations with the user's analytic task. It allows users to input their tasks and interests flexibly, thereby striking a better balance between automation and the user's intent. We hope that it can inspire more valuable works in the VisRec community. Nevertheless, it is still in the early stage of visualization recommendation. We now discuss how TaskVis can evolve to be more intelligent and what we need to do in the future.

**Extract abundant data properties:** The rules in TaskVis are more detailed to the task but less involved in data properties, where only type, length, and cardinality of data columns are considered. However, the properties of multi-dimensional statistical data are beneficial for targeted recommendations, such as standard deviation, coefficient of variance, and quantitative coefficient of dispersion that extracted in VizML [25]. In the future, we will extend the rule base to support more statistical data properties.

**Integrate semantics and natural language interface:** TaskVis only integrates data properties in the numerical sense but does not consider the semantic meaning. In the future, we will integrate the semantic type of data columns into TaskVis to transmit more valuable information [28]. Further, the user's occupation, social hotspots, and geographic information can be helpful to the personalized recommendation. In addition, a natural language interface can also greatly improve the user experience of the VisRec systems [47].

**Knowledge graph to construct rules:** Draco [37] has formalized visualization design knowledge as constraints, but the knowledge is independent of each other, which is not conducive to knowledge reasoning. We suppose leveraging a knowledge graph to construct a rule base in an extensible model where the rules are interrelated. Based on the Knowledge graph, we can conduct knowledge extraction from public examples. In addition, users can also be allowed for customizing rules to satisfy their requirements.



## REFERENCES

- [1] Cars dataset. <https://vega.github.io/editor/data/cars.json>.
- [2] Chart Chooser. <http://labs.juiceanalytics.com/chartchooser/>.
- [3] ChartGuide poster. <https://chart.guide/topics/chartguide-poster-4-0/>.
- [4] Choosing a good chart. <https://extremepresentation.typepad.com/blog/>.
- [5] COVID-19 dataset. <https://github.com/CSSEGISandData/COVID-19>.
- [6] Data Visualization Catalogue. <https://datavizcatalogue.com/index.html>.
- [7] Hollywood stories dataset. <https://www.kaggle.com/brendan45774/hollywood-most-profitable-stories>.
- [8] System Usability Scale. <https://www.usability.gov/how-to-and-tools/methods/system-usability-scale.html>.
- [9] The Graphic Continuum. <https://policyviz.com/2014/09/09/graphic-continuum/>.
- [10] Visual Analytics. <https://www.pinterest.com/pin/20125529565819990/>.
- [11] Visual Vocabulary. <http://ft-interactive.github.io/visual-vocabulary/>.
- [12] World happiness dataset. <https://www.promptcloud.com/world-happiness-report-dataset-2019/>.
- [13] R. Amar, J. Eagan, and J. Stasko. Low-level components of analytic activity in information visualization. In *Proc. INFOVIS'05*. IEEE, 2005.
- [14] M. Brehmer and T. Munzner. A Multi-Level Typology of Abstract Visualization Tasks. *IEEE Trans. Vis. Comput. Graph.*, 19(12), 2013.
- [15] M. Correll and M. Gleicher. Error Bars Considered Harmful: Exploring Alternate Encodings for Mean and Error. *IEEE Trans. Vis. Comput. Graph.*, 20(12), 2014.
- [16] F. E. Croxton and R. E. Stryker. Bar Charts Versus Circle Diagrams. *J. Am. Stat. Assoc.*, 22(160):473, 1927.
- [17] Z. Cui, S. K. Badam, M. A. Yalçin, and N. Elmqvist. DataSite: Proactive visual data exploration with computation of insight-based recommendations. *Inf. Vis.*, 18(2), 2019.
- [18] Ç. Demiralp, P. J. Haas, S. Parthasarathy, and T. Pedapati. Foresight: Recommending visual insights. *Proc. VLDB Endow.*, 10(12):1937–1940, 2017.
- [19] C. D. Demiralp, M. S. Bernstein, and J. Heer. Learning Perceptual Kernels for Visualization Design. *IEEE Trans. Vis. Comput. Graph.*, 20(12):1933–1942, 2014.
- [20] V. Dibia and C. Demiralp. Data2Vis: Automatic Generation of Data Visualizations Using Sequence-to-Sequence Recurrent Neural Networks. *IEEE Comput. Graph. Appl.*, 39(5):33–46, 2019.
- [21] W. C. Eells. The Relative Merits of Circles and Bars for Representing Component Parts. *J. Am. Stat. Assoc.*, 21(154):119, 1926.
- [22] M. Gebser, B. Kaufmann, R. Kaminski, and et al. Potassco: The Potsdam Answer Set Solving Collection. *AI Commun.*, 24(2):107–124, 2011.
- [23] D. Gotz and Z. Wen. Behavior-driven visualization recommendation. In *Proc. IUI'09*. ACM, 2009.
- [24] C. Harris, R. A. Rossi, S. Malik, and et al. Insight-centric Visualization Recommendation. *arXiv*, 2021.
- [25] K. Hu, M. A. Bakker, S. Li, and et al. VizML: A machine learning approach to visualization recommendation. In *Proc. CHI'19*. ACM, 2019.
- [26] K. Hu, S. S. Gaikwad, M. Hulsebos, and et al. VizNet: Towards a large-scale visualization learning and benchmarking repository. In *Proc. CHI'19*. ACM, 2019.
- [27] K. Hu, D. Orghian, and C. Hidalgo. DIVE: A mixed-initiative system supporting integrated data exploration workflows. In *Proc. HILDA'2018*. ACM, 2018.
- [28] M. Hulsebos, A. Satyanarayan, K. Hu, and et al. Sherlock: A deep learning approach to semantic data type detection. In *Proc. KDD'19*. ACM, 2019.
- [29] S. Kandel, R. Parikh, A. Paepcke, and et al. Profiler: Integrated statistical analysis and visualization for data quality assessment. In *Proc. AVI'12*. ACM, 2012.
- [30] S. Kerpedjiev, G. Carenini, S. F. Roth, and J. D. Moore. AutoBrief: a multimedia presentation system for assisting data analysis. *Comput. Stand. Interfaces*, 18(6-7), 1997.
- [31] Y. Kim and J. Heer. Assessing Effects of Task and Data Distribution on the Effectiveness of Visual Encodings. *Comput. Graph. Forum*, 37(3):157–167, 2018.
- [32] Y. Kim, K. Wongsuphasawat, J. Hullman, and J. Heer. GraphScape: A model for automated reasoning about visualization similarity and sequencing. In *Proc. CHI'17*. ACM, 2017.
- [33] H. Lin, D. Moritz, and J. Heer. Dziban: Balancing Agency & Automation in Visualization Design via Anchored Recommendations. In *Proc. CHI'20*. ACM, 2020.
- [34] Y. Luo, X. Qin, N. Tang, and G. Li. Deepeye: towards automatic data visualization. In *Proc. ICDE'18*. IEEE.
- [35] J. Mackinlay. Automating the design of graphical presentations of relational information. *ACM Trans. Graph.*, 5(2):110–141, 1986.
- [36] J. Mackinlay, P. Hanrahan, and C. Stolte. Show Me: Automatic Presentation for Visual Analysis. *IEEE Trans. Vis. Comput. Graph.*, 13(6):1137–1144, 2007.
- [37] D. Moritz, C. Wang, G. L. Nelson, and et al. Formalizing Visualization Design Knowledge as Constraints: Actionable and Extensible Models in Draco. *IEEE Trans. Vis. Comput. Graph.*, 25(1):438–448, 2019.
- [38] A. Narechania, A. Srinivasan, and J. Stasko. NL4DV: A Toolkit for Generating Analytic Specifications for Data Visualization from Natural Language Queries. *IEEE Trans. Vis. Comput. Graph.*, 27(2):369–379, 2021.
- [39] X. Qian, R. A. Rossi, F. Du, and et al. Learning to Recommend Visualizations from Data. In *Proc. KDD'21*. ACM, 2021.
- [40] X. Qin, Y. Luo, N. Tang, and G. Li. Making data visualization more efficient and effective: a survey. *VLDB J.*, 29(1):93–117, 2020.
- [41] A. Rind, W. Aigner, M. Wagner, and et al. Task Cube: A three-dimensional conceptual space of user tasks in visualization design and evaluation. *Inf. Vis.*, 15(4), 2016.
- [42] S. F. Roth, J. Kolojechick, J. Mattis, and J. Goldstein. Interactive graphic design using automatic presentation knowledge. In *Proc. CHI'94*. ACM, 1994.
- [43] B. Saket, A. Endert, and C. Demiralp. Task-Based Effectiveness of Basic Visualizations. *IEEE Trans. Vis. Comput. Graph.*, 25(7):2505–2512, 2019.
- [44] A. Sarikaya and M. Gleicher. Scatterplots: Tasks, Data, and Designs. *IEEE Trans. Vis. Comput. Graph.*, 24(1):402–412, 2018.
- [45] A. Satyanarayan, D. Moritz, K. Wongsuphasawat, and J. Heer. Vega-Lite: A Grammar of Interactive Graphics. *IEEE Trans. Vis. Comput. Graph.*, 23(1):341–350, 2017.
- [46] J. Seo and B. Shneiderman. A Rank-by-Feature Framework for Interactive Exploration of Multidimensional Data. *Inf. Vis.*, 4(2), 2005.
- [47] L. Shen, E. Shen, Y. Luo, and et al. Towards Natural Language Interfaces for Data Visualization: A Survey. *IEEE Trans. Vis. Comput. Graph.*, pp. 1–20, 2022.
- [48] D. Shi, Y. Shi, X. Xu, and et al. Task-Oriented Optimal Sequencing of Visualization Charts. In *Proc. VDS'19*. IEEE, 2019.
- [49] D. Skau and R. Kosara. Arcs, Angles, or Areas: Individual Data Encodings in Pie and Donut Charts. *Comput. Graph. Forum*, 2016.
- [50] A. Srinivasan, S. M. Drucker, A. Endert, and J. Stasko. Augmenting Visualizations with Interactive Data Facts to Facilitate Interpretation and Communication. *IEEE Trans. Vis. Comput. Graph.*, 25(1), 2019.
- [51] B. Steichen, G. Carenini, and C. Conati. User-adaptive information visualization - Using eye gaze data to infer visualization tasks and user cognitive abilities. In *Proc. IUI'13*. ACM, 2013.
- [52] C. Stolte, D. Tang, and P. Hanrahan. Polaris: a system for query, analysis, and visualization of multidimensional relational databases. *IEEE Trans. Vis. Comput. Graph.*, 8(1):52–65, 2002.
- [53] M. Vartak, S. Huang, T. Siddiqui, and et al. Towards Visualization Recommendation Systems. *ACM SIGMOD Rec.*, 45(4):34–39, 2017.
- [54] M. Vartak, S. Rahman, S. Madden, and et al. SEEDB: Efficient data-driven visualization recommendations to support visual analytics. *Proc. VLDB Endow.*, 8(13), 2015.
- [55] Y. Wang, F. Han, L. Zhu, and et al. Line Graph or Scatter Plot? Automatic Selection of Methods for Visualizing Trends in Time Series. *IEEE Trans. Vis. Comput. Graph.*, 24(2):1141–1154, 2018.
- [56] Y. Wang, Z. Sun, H. Zhang, and et al. DataShot: Automatic Generation of Fact Sheets from Tabular Data. *IEEE Trans. Vis. Comput. Graph.*, 26(1):895–905, 2020.
- [57] Y. Wang, Z. Sun, H. Zhang, and et al. DataShot: Automatic Generation

- of Fact Sheets from Tabular Data. *IEEE Trans. Vis. Comput. Graph.*, 26(1):895–905, 2020.
- [58] K. Wongsuphasawat, D. Moritz, A. Anand, and et al. Towards a general-purpose query language for visualization recommendation. In *Proc. HILDA'16*. ACM, 2016.
- [59] K. Wongsuphasawat, Z. Qu, D. Moritz, and et al. Voyager 2: Augmenting visual analysis with partial view specifications. In *Proc. CHI'17*. ACM, 2017.
- [60] A. Wu, Y. Wang, M. Zhou, and et al. MultiVision: Designing Analytical Dashboards with Deep Learning Based Recommendation. *IEEE Trans. Vis. Comput. Graph.*, pp. 1–11, 2021.
- [61] M. Zhou, Q. Li, X. He, and et al. Table2Charts: Recommending Charts by Learning Shared Table Representations. In *Proc. KDD'21*. ACM, 2021.
- [62] S. Zhu, G. Sun, Q. Jiang, and et al. A survey on automatic infographics and visualization recommendations. *Vis. Informatics*, 4(3), 2020.